

Zero Slop

You used AI to help you write, and now it sounds like AI. Zero Slop finds that sound and takes it out, without changing what you said.

```
zero-slop - before / after

$ pbpaste | slopscore --explain
"We're thrilled to announce that our team has leveraged cutting-edge
AI to deliver a seamless onboarding experience, cutting setup 40%."

AI-likelihood: 100.0/100 [slop]
charged spans (6), heaviest first:
  5 linkedin    announce-excited  "We're thrilled to"
  4 marketing   marketing-register 'cutting-edge'
  4 marketing   marketing-register 'seamless'
  4 lexicon     cutting-edge    'cutting-edge'
  4 lexicon     seamless      'seamless'
  2 rider      leverage       'leveraged'

$ deslop | slopscore --explain
"We cut onboarding setup time by 40% using AI."

AI-likelihood: 9.5/100 [clean]
charged spans: none - the score is rhythm and format only
```

What it does

A linter and rewriter for people who publish under their own name. It tells you whether a reader will think a machine wrote your draft, and it removes the parts that give that impression, without touching your facts. It is a de-slopper, not a writer: it subtracts the tells and returns your words, structure and voice intact. It is **not** a detector, and it does not optimise against detector scores — a losing game now that adversarial paraphrasing cuts detector accuracy 87–99% (arXiv:2506.07001).

How it decides

Slop is not one thing, so it looks at four, and only the first pays attention to the actual words. That is the point: you cannot dodge it with a thesaurus, because most of what it measures is the shape of the writing. It leans toward mercy, too. One em-dash is not slop; plenty of great writing uses them. A tell only counts once the word choice backs it up, so a cluster convicts and a lone one does not.

Does it actually work

Method	Best-picks	Composite
Zero Slop	55	8.01
blader/humanizer	40	7.82
petergyang/no-ai-slop	5	6.96
isatimur/de-slop	0	6.35

We tested it against the four tools it is built on, on 50 AI-heavy drafts judged blind. Across 100 picks, judges chose the Zero Slop version most often, 55 to blader's 40. Honestly: it clearly beat two of the four and tied the strongest. Every accuracy number here comes from test writing we made ourselves, so it shows the tool agreeing with an obvious call, not that it will nail every draft in the wild.

It will not touch your facts

Most 'humanizer' tools will invent a detail or drop a number to make a sentence flow. Zero Slop treats that as the one unforgivable move, and --fidelity checks every figure, name, quote and link survived and nothing new snuck in. We built it because an early version of the tool did exactly that in testing, inventing a feeling the writer never mentioned. We would rather show you the miss than hide it.

It gets sharper as you use it

Most tools are frozen the day they ship. This one learns from you. What you delete after it hands your draft back was a tell it missed, and a phrase only becomes a rule after three different documents cut it, so one person's quirk never becomes everyone's rule. A sample of your own writing teaches it which words are just how you talk.

MIT · github.com/manavmishra/ZeroSlop · v1.6.0 · 62 tests · offline, zero dependencies