

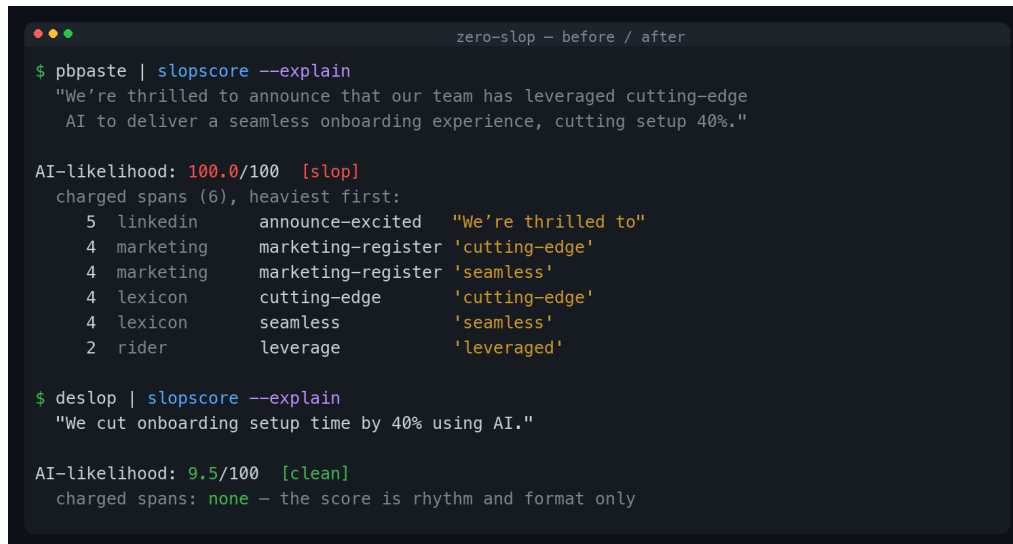
Zero Slop

You used AI to help with your writing, and now it reads like a machine wrote every word. You can hear it, and so can everyone who reads it. There is a word for that machine sound now: slop.

Zero Slop finds that slop in your draft and takes it out, without changing what you actually said. It scores the writing from 0 to 100, points at the exact words dragging it down, rewrites them, and then double-checks that it kept every fact you had and invented nothing.

```
npx skills add manavmishra/ZeroSlop --global
```

Then say "de-slop this" in any agent. No account, no server, and your writing never leaves your machine.



```
zero-slop - before / after

$ pbpaste | slopscore --explain
"We're thrilled to announce that our team has leveraged cutting-edge
AI to deliver a seamless onboarding experience, cutting setup 40%."

AI-likelihood: 100.0/100 [slop]
charged spans (6), heaviest first:
  5 linkedin    announce-excited "We're thrilled to"
  4 marketing   marketing-register 'cutting-edge'
  4 marketing   marketing-register 'seamless'
  4 lexicon     cutting-edge   'cutting-edge'
  4 lexicon     seamless      'seamless'
  2 rider       leverage     'leveraged'

$ deslop | slopscore --explain
"We cut onboarding setup time by 40% using AI."

AI-likelihood: 9.5/100 [clean]
charged spans: none - the score is rhythm and format only
```

Why it matters now

Sounding like AI used to just cost you a little credibility. Now it costs you reach. Over two weeks in mid-2026, LinkedIn, Snapchat, YouTube, and Substack all started acting on slop. On LinkedIn, a flagged post now stops reaching anyone outside your own network. One study the platforms keep citing estimates that up to 40% of social-media writing is already machine-made. When a feed is that full of slop, the systems ranking it bury anything that smells the same, and human-sounding writing is what gets through.

Who it is for

Anyone who writes under their own name and would rather not sound like a robot. A founder writing a launch post, a marketer shipping five things a week, a researcher who needs their paper to stay formal, an engineering team that wants slop to fail a check the way a bug does.

How it decides

First, it measures. Zero Slop scores your draft from 0 to 100, and every point of that score points back at a phrase you can see, so the number is never a black box. Slop is not one thing, so it weighs several signals, and only the first pays attention to your actual words. That is the point: you cannot dodge it with a thesaurus, because most of what it measures is the shape of the writing, not the vocabulary. It weighs the words you chose, but also the rhythm of your sentences, how hard they are to read, and how heavily the page is formatted. Then it leans toward mercy. One em-dash is not slop; plenty of great writing leans on them. A tell only counts once the word choice backs it up, so a cluster convicts and a lone one does not. It adds a check no rulebook can: it asks the model running it to guess your words, and counts how often it can — machine writing is easy to predict, and yours should not be. Zero

Slop ships no model of its own; it borrows whichever one you are already in. When it rewrites, it writes several versions and keeps the cleanest one that changes no fact.

What makes it different

It is built on four open-source projects that worked out this craft first: no-ai-slop, humanizer, de-slop, and stop-slop. It adds three things a plain rewriter cannot.

A score you can trust and check. It comes from rules you can read, not a black box, so every point traces to a phrase you can see, and it can go straight into a build.

A promise about your facts. Most "humanizer" tools will invent a detail or drop a number to make a sentence flow. Zero Slop lists every figure, name, quote, and link in your original and fails a rewrite that loses one or adds one. A missing number you would catch; an invented one reads perfectly and slips right past, which is exactly why the check exists.

A tool that learns you. This is the part most tools skip. After it hands your draft back and you edit it, it compares the two: a phrase you cut was a tell it should have caught, and one it flagged that you kept was a false alarm. A phrase only becomes a new rule after three separate people's drafts have cut it, so no single quirk speaks for everyone — and the meter gets sharper the more the tool is used, instead of staying frozen. A sample of your own writing also teaches it which words are just how you talk.

The tool even improves itself. The rewrite follows a page of written instructions, and those can be tuned like anything with a score. Microsoft's SkillOpt edits them, runs the edited version on a batch of drafts, and grades each result the way your writing is graded — how much slop came out, and whether every fact stayed in. It keeps an edit only when it beats the old instructions on drafts it has never seen, and discards the rest. Every accepted edit makes the next de-slop a little sharper, for everyone who uses the skill, not just one draft. It also keeps itself current: each session it checks for a newer release and points you at the one-line update, sending a version query and nothing of what you wrote.

Does it actually work

We ran it head to head against the four tools it builds on: 50 AI-heavy drafts, six kinds of writing, judged blind. Across 100 picks, judges chose the Zero Slop version most often, 55 to humanizer's 40, with the other two far back. It clearly beat two of them, and tied the strongest. It also cleanly separates obvious slop from genuine human writing, with no overlap, at about 1,100 documents a second.

The honest caveat, which the full README does not bury: our accuracy numbers all come from test writing we created ourselves, so they show the tool agreeing with an obvious call, not that it will nail every draft in the wild. A number you could really trust needs about a thousand hand-labeled samples, and building that set is the next thing on the list.