

Zero Slop

Add taste to your output, or fix your slop. It scores a draft 0–100, shows which words cost what, rewrites it, and checks the rewrite kept every fact and made nothing up.

```
zero-slop - before / after

$ pbpaste | slopscore --explain
"We're thrilled to announce that our team has leveraged cutting-edge
AI to deliver a seamless onboarding experience, cutting setup 40%."

AI-likelihood: 100.0/100 [slop]
charged spans (6), heaviest first:
  5 linkedin announce-excited "We're thrilled to"
  4 marketing marketing-register 'cutting-edge'
  4 marketing marketing-register 'seamless'
  4 lexicon cutting-edge 'cutting-edge'
  4 lexicon seamless 'seamless'
  2 rider leverage 'leveraged'

$ deslop | slopscore --explain
"We cut onboarding setup time by 40% using AI."

AI-likelihood: 9.5/100 [clean]
charged spans: none - the score is rhythm and format only
```

What it is, and is not

It is a linter and rewriter for people who publish under their own name. It tells you whether a reader will think a machine wrote your draft, and it fixes the parts that give them that impression, without ever touching your facts. What it is **not** is a detector. It will not classify anyone's work as AI-generated, and it will not optimise against detector scores — a losing game now that adversarial paraphrasing cuts detector accuracy 87–99% (arXiv:2506.07001).

How it works

Four channels score every draft, and three never look at specific words, so a synonym swap does not move the number. They fuse cautiously: a lone stylistic habit like an em-dash describes plenty of excellent human writing, so it only counts once the vocabulary channel agrees. Clusters convict, singles do not.

The numbers

Method	Best-picks	Composite
Zero Slop	55	8.01
blader/humanizer	40	7.82
petergyang/no-ai-slop	5	6.96
isatimur/de-slop	0	6.35

Pooled over 100 blind verdicts on 50 drafts across six genres. It wins the plurality against a 25% chance rate ($p = 1.7e-10$), and it is statistically tied with blader head to head ($p = 0.15$). Telling obvious slop from obvious human writing across five genres: AUC 1.000, 77-point separation. False positives on ordinary human prose: zero of five. Throughput about 1,100 docs/sec.

The honest weakness

In the benchmark, Zero Slop came **last of four on judge-rated fidelity**, and it was the only tool to carry a fabrication flag, for a rewrite that handed the author a feeling they had never claimed. Nothing in the loop measured fidelity at the time, so the property the tool cares about most was riding on instruction alone. That is the whole reason --fidelity exists, and it now catches the mechanical half, an invented figure or name or feeling, while a claim quietly reframed still needs the judgment pass. Publishing the result that went against it is the reason to trust the ones that did not.

It learns from your edits, and adapts to you

A span becomes a new pattern only once three separate documents have cut it, and a pattern you overrule three times loses half its weight, so the meter sharpens in both directions the more it is used. It will also learn you in particular: build a profile from a sample of your own writing, and it stops charging you for the tell-words you reach for by reflex, while everyone else's meter stays untouched and real slop still gets caught.

MIT · github.com/manavmishra/ZeroSlop · v1.6.0 · 62 tests · offline, zero dependencies