


Zero Slop

A linter for the AI accent. It scores a draft 0–100, shows you which spans cost what, rewrites it, and checks the rewrite kept every fact and made nothing up.

What it is, and is not

It is a linter and rewriter for people who publish under their own name. It tells you whether a reader will think a machine wrote your draft, and it fixes the parts that give them that impression, without ever touching your facts. What it is **not** is a detector. It will not classify anyone's work as AI-generated, and it will not optimise against detector scores — a losing game now that adversarial paraphrasing cuts detector accuracy 87–99% (arXiv:2506.07001).

How it works — four interpretable channels



Three of the four channels never look at specific words at all, which is why a score holds up even when someone swaps synonyms to slip past a word list. The fusion is cautious by design, because a single stylistic habit is not evidence of anything. An em-dash or a formal register describes a great deal of excellent human writing, so those only carry weight once the lexical channel backs them up. Clusters convict, singles do not.

The numbers

Method	Best-picks	Composite
Zero Slop	55	8.01
blader/humanizer	40	7.82
petergyang/no-ai-slop	5	6.96
isatimur/de-slop	0	6.35

Pooled over 100 blind verdicts on 50 drafts across six genres. It wins the plurality against a 25% chance rate ($p = 1.7e-10$), and it is statistically tied with blader head to head ($p = 0.15$). Telling obvious slop from obvious human writing across five genres: AUC 1.000, 77-point separation. False positives on ordinary human prose: zero of five. Throughput about 1,100 docs/sec.

The honest weakness

In the benchmark, Zero Slop came **last of four on judge-rated fidelity**, and it was the only tool to carry a fabrication flag, for a rewrite that handed the author a feeling they had never claimed. Nothing in the loop measured fidelity at the time, so the property the tool cares about most was riding on instruction alone. That is the whole reason --fidelity exists, and it now catches the mechanical half, an invented figure or name or feeling, while a claim quietly reframed still needs the judgment pass. Publishing the result that went against it is the reason to trust the ones that did not.

It learns from your edits, and adapts to you

A span becomes a new pattern only once three separate documents have cut it, and a pattern you overrule three times loses half its weight, so the meter sharpens in both directions the more it is used. It will also learn you in particular: build a profile from a sample of your own writing, and it stops charging you for the tell-words you reach for by reflex, while everyone else's meter stays untouched and real slop still gets caught.

MIT · github.com/manavmishra/ZeroSlop · v1.6.0 · 62 tests · offline, zero dependencies