

Zero Slop

A linter for the AI accent — scores a draft 0–100, shows which spans cost what, rewrites it, and verifies the rewrite invented nothing.

What it is, and is not

A linter and rewriter for people publishing under their own name. It tells you whether a reader will think a machine wrote your draft, and fixes the parts that make them think so, without touching your facts. It is **not** a detector: it will not classify anyone's work as AI-generated, and it deliberately does not optimise against detector scores — a losing game, since 2026 research shows adversarial paraphrasing cuts detector accuracy 87–99% (arXiv:2506.07001).

How it works — four interpretable channels



Three of the four never look at specific words, which is why a score survives someone swapping synonyms to dodge a word list. The fusion is cautious: a single stylistic habit is not evidence, so em-dashes and formal register only carry weight when the lexical channel corroborates. Clusters convict, singles do not.

The numbers

Method	Best-picks	Composite
Zero Slop	55	8.01
blader/humanizer	40	7.82
petergyang/no-ai-slop	5	6.96
isatimur/de-slop	0	6.35

Pooled over 100 blind verdicts on 50 drafts, six genres. Wins the plurality against a 25% chance rate ($p = 1.7\text{e-}10$); statistically tied with blader head to head ($p = 0.15$). Discrimination between obvious slop and obvious human writing across five genres: AUC 1.000, 77-point separation. False positives on ordinary human prose: 0 of 5. Throughput ~1,100 docs/sec.

The honest weakness

In the benchmark, Zero Slop ranked **last of four on judge-rated fidelity** and carried the only fabrication flag — a rewrite invented a feeling the author never described. Nothing in the loop measured fidelity, so the property it claims to protect was enforced by instruction alone. The --fidelity channel exists because of that and now catches invented figures, names and stated feelings. Publishing the result that went against it is the reason to trust the ones that did not. Every accuracy number rests on corpora the author wrote; a defensible rate needs a labelled corpus at volume, which is the top roadmap item.

It learns from you, and adapts to you

A span becomes a new pattern only after three independent documents cut it; a pattern you overrule three times loses half its weight, so the meter sharpens in both directions with use. And it personalises: build a profile from a sample of your own writing, and the meter stops charging you for the tell-words you genuinely use, while everyone else's meter is unchanged and real slop is still caught.

MIT · github.com/manavmishra/ZeroSlop · v1.6.0 · 62 tests · offline, zero dependencies